

speedup over torch.compile

- TF block (n=256k)
- reduction (n=256k)
- attention (n=256k)

20x
15x
10x
5x
0x

float64

float32

float16

