




steady-state time per iteration (μs, log scale)

lower is better

MetaTensor

torch.compile (max-autotune)

JAX/XLA

torch.compile

PyTorch eager

Torch-TensorRT

torch.compile (reduce-overhead)