





kernel launches per forward, lower is better



GPU busy / wall time, higher is better

MetaTensor torch.compile JAX/XLA