



next-token latency, ms (lower is better)

- MetaTensor
- PyTorch eager
- torch.compile
- torch.compile (reduce-overhead)